



# AutoMLOps

*From Notebooks to Pipelines in Minutes*

**Implementation Guide v1.0.1**

srastatter@, tonydiloireto@, allegranoto@

February 2023

The diagram illustrates the AutoMLOps CI/CD & Orchestration architecture on Google Cloud Platform. The workflow is as follows:

- Call AutoMLOps** (Vertex AI Workbench) triggers the **Generate/Run Code** (AutoMLOps.go) and **Generate Code** (AutoMLOps.generate) processes.
- Generate/Run Code** and **Generate Code** both trigger the **Source Code** (CSR) process.
- Source Code** (CSR) triggers the **Run Build on Push** (Cloud Build Trigger) process.
- Run Build on Push** (Cloud Build Trigger) triggers three parallel processes:
  - Build Base Image** (Cloud Build)
  - Build Runner SVC** (Cloud Build)
  - Deploy Runner SVC** (Cloud Build)
- Build Base Image** (Cloud Build) and **Build Runner SVC** (Cloud Build) both trigger the **Packages** (Artifact Registry) process.
- Packages** (Artifact Registry) triggers the **Run PipelineJob** (Cloud Run) process.
- Deploy Runner SVC** (Cloud Build) triggers the **Run PipelineJob** (Cloud Run) process.
- Run PipelineJob** (Cloud Run) triggers the **Training Job** (Vertex AI Pipelines) process.
- Job Queueing SVC** (Cloud Tasks) triggers the **Run PipelineJob** (Cloud Run) process, with a label **<Invokes>** indicating the relationship.
- Scheduled Run** (Cloud Scheduler) triggers the **Run PipelineJob** (Cloud Run) process.

The final output of the pipeline is the **Training Job** (Vertex AI Pipelines).

Set Up



# Prerequisites / Assumptions

*The prerequisites for use of AutoMLOps are as follows:*

- Jupyter (or Jupyter-compatible) notebook environment
- [Notebooks API](#) is enabled
- Python version  $\geq 3.0$  and  $\leq 3.10$
- [Google Cloud SDK 407.0.0](#)
- [gcloud beta 2022.10.21](#)
- [Terraform](#) is installed
- git is installed and logged-in

```
git config --global user.email "you@example.com"  
git config --global user.name "Your Name"
```

- [Application Default Credentials \(ADC\)](#) are set up, which can be done through the following commands:

```
gcloud auth application-default login  
gcloud config set account <account@example.com>
```

# Set Up AutoMLOps Package

1. Clone the repo:

```
git clone https://github.com/GoogleCloudPlatform/AutoMLOps.git
```

2. Install the AutoMLOps package:

```
pip install dist/AutoMLOps-1.0.1-py2.py3-none-any.whl
```

3. Open a notebook and import the AutoMLOps package:

```
from AutoMLOps import AutoMLOps
```

4. Decide whether to use Kubeflow definitions or Python definitions

## Using Python Components (no Kubeflow)



# AutoMLOps Python Process

1. Define import code cells

```
%%define_imports
import json
import pandas as pd
from google.cloud import aiplatform
from google.cloud import aiplatform_v1
from google.cloud import bigquery
...
```

# AutoMLOps Python Process

## 2. Define component code cells

```
%%define_component
AutoMLOps.makeComponent(
    name = "create_dataset",
    description = "Loads data from BQ and writes a dataframe as a csv to GCS.", # optional
    params = [
        {"name": "bq_table", "type": str, # descriptions are optional
        {"name": "data_path", "type": str, "description": "GS location where the training data is written."},
        {"name": "project_id", "type": str, "description": "Project_id."}
    ]
)
```



# AutoMLOps Python Process

## 3. Define pipeline code cell

```
AutoMLOps.makePipeline(  
    name = "training-pipeline",  
    description = "description", # optional  
    params = [  
        {"name": "bq_table", "type": str}, # descriptions are optional  
        {"name": "model_directory", "type": str, "description": "Description."},  
        {"name": "data_path", "type": str, "description": "Description."},  
        {"name": "project_id", "type": str, "description": "Description."},  
        {"name": "region", "type": str, "description": "Description."}  
    ],  
    pipeline = [{  
        "component_name": "create_dataset", "param_mapping": [  
            ("bq_table", "bq_table"), # (component_param, pipeline_param)  
            ("data_path", "data_path"),  
            ("project_id", "project_id")  
        ]  
    },  
    {  
        "component_name": "train_model", "param_mapping": [  
            ("model_directory", "model_directory"),  
            ("data_path", "data_path")  
        ]  
    },  
    {  
        "component_name": "deploy_model", "param_mapping": [  
            ("model_directory", "model_directory"),  
            ("project_id", "project_id"),  
            ("region", "region")  
        ]  
    }  
    ]  
)
```

# AutoMLOps Python Process

4. Define the pipeline parameters dictionary

```
pipeline_params = {  
    "bq_table": f"{PROJECT_ID}.test_dataset.dry-beans",  
    "model_directory": f"gs://{PROJECT_ID}-bucket/trained_models/{datetime.datetime.now()}",  
    "data_path": f"gs://{PROJECT_ID}-bucket/data",  
    "project_id": f"{PROJECT_ID}",  
    "region": "us-central1"  
}
```

# AutoMLOps Python Process

5. Call `AutoMLOps.generate()` to create the resources and repository  
Or `AutoMLOps.go()` to call generate in addition to building/submitting the pipeline job

```
AutoMLOps.generate(project_id = PROJECT_ID,  
                    pipeline_params = pipeline_params,  
                    use_kfp_spec = False,  
                    run_local = True)  
  
AutoMLOps.go(project_id = PROJECT_ID,  
              pipeline_params = pipeline_params,  
              use_kfp_spec = False,  
              run_local = True)
```

Set `use_kfp_spec=False` when using an AutoMLOps Python defined pipeline

Set `run_local=False` if you want to generate and use CI/CD features

## Using Kubeflow Components



# AutoMLOps Kubeflow Process

1. Define your components using KFP

```
@component(  
    packages_to_install = [  
        "google-cloud-bigquery",  
        "pandas",  
        "pyarrow",  
        "db_dtypes"  
    ],  
    base_image = "python:3.9",  
    output_component_file = f"{AutoMLOps.OUTPUT_DIR}/create_dataset.yaml"  
)  
def create_dataset(  
    bq_table: str,  
    output_data_path: OutputPath("Dataset"),  
    project: str  
):  
    from google.cloud import bigquery  
    ...
```

# AutoMLOps Kubeflow Process

2. If using Kubeflow defs, define your pipeline using KFP and AutoMLOps:

```
%%define_kfp_pipeline

@dsl.pipeline(name = 'training-pipeline')
def pipeline(bq_table: str,
             output_model_directory: str,
             project: str,
             region: str,
             ):

    dataset_task = create_dataset(
        bq_table = bq_table,
        project = project)

    model_task = train_model(
        output_model_directory = output_model_directory,
        dataset = dataset_task.output)

    deploy_task = deploy_model(
        model = model_task.outputs["model"],
        project = project,
        region = region)

    ...
```

# AutoMLOps Kubeflow Process

## 3. Define the pipeline parameters dictionary

```
pipeline_params = {  
    "bq_table": f"{PROJECT_ID}.test_dataset.dry-beans",  
    "output_model_directory": f"gs://{PROJECT_ID}-bucket/trained_models/{datetime.datetime.now()}",  
    "project": f"{PROJECT_ID}",  
    "region": "us-central1"  
}
```

# AutoMLOps Kubeflow Process

4. Call `AutoMLOps.generate()` to create the resources and repository  
Or `AutoMLOps.go()` to call generate in addition to building/submitting the pipeline job

```
AutoMLOps.generate(project_id = PROJECT_ID,  
                    pipeline_params = pipeline_params,  
                    use_kfp_spec = True,  
                    run_local = False)  
  
AutoMLOps.go(project_id = PROJECT_ID,  
              pipeline_params = pipeline_params,  
              use_kfp_spec = True,  
              run_local = False)
```

Set `use_kfp_spec=True` when using a Kubeflow defined pipeline

Set `run_local=False` if you want to generate and use CI/CD features



## Behind the Scenes



# Cloud Resources

When the pipeline is run, the following resources are created to complete the MLOps pipeline:

1. AutoMLOps codebase
2. Artifact Registry
3. GS Bucket
4. Pipeline Runner Service Account
5. Cloud Source Repository (turns the notebooks working directory into a CSR)
6. Cloud Build Trigger
7. Cloud Runner Service
8. Cloud Scheduler
9. Cloud Tasks queue

# APIs

When the pipeline is run, the following APIs are enabled:

1. `cloudresourcemanager.googleapis.com`
2. `aiplatform.googleapis.com`
3. `artifactregistry.googleapis.com`
4. `cloudbuild.googleapis.com`
5. `cloudscheduler.googleapis.com`
6. `cloudtasks.googleapis.com`
7. `compute.googleapis.com`
8. `iam.googleapis.com`
9. `iamcredentials.googleapis.com`
10. `ml.googleapis.com`
11. `run.googleapis.com`
12. `storage.googleapis.com`
13. `sourcerepo.googleapis.com`

# IAM Access

When the pipeline is run, the following IAM access roles are updated:

1. **Pipeline Runner Service Account** (created if it does exist, defaults to: *vertex-pipelines@<PROJECT\_ID>.iam.gserviceaccount.com*).

Roles added:

- roles/aiplatform.user
- roles/artifactregistry.reader
- roles/bigquery.user
- roles/bigquery.dataEditor
- roles/iam.serviceAccountUser
- roles/storage.admin
- roles/run.admin

2. **Cloudbuild Default Service Account** (<PROJECT\_NUMBER>@cloudbuild.gserviceaccount.com).

Roles added:

- roles/run.admin
- roles/iam.serviceAccountUser
- roles/cloudtasks.enqueueur
- roles/cloudscheduler.admin

# Package Dependencies

When using AutoMLOps, the following package versions are used:

1. autoflake==2.0.0,
2. docopt==0.6.2,
3. ipython==7.34.0,
4. pipreqs==0.4.11,
5. pyflakes==3.0.1,
6. PyYAML==5.4.1,
7. yarg==0.1.9